

Organized by:




اونيفورسيتي تيكنولوگي بروني
UNIVERSITI TEKNOLOGI BRUNEI

Supported by:



Sponsored by:



INNOVATE FOR IMPACT SDG Hackathon 2025

Are you aged 18 to 30 and passionate about solving real-world problems with technology? Join the AI & IoT for SDG Hackathon, where your ideas can drive meaningful impact!

Use Artificial Intelligence or Internet of Things (IoT) to develop innovative solutions aligned with any of the 17 SDGs. Whether it's a simulation, prototype, or working demo, we want to see your tech-driven vision for a better world.

Top Prize: **BND 1,200** (approx. EUR 900) subject to exchange rate at time of award.

Free registration! Sign up by June 15th. A registration number and submission link will be provided upon signing up.

Submission Requirements:

- Work individually or in teams.
- Submit a 5-minute pitch video.
- Submit a solution report (max. 10 pages) explaining your concept, design, impact, and implementation plan.
- Submission period: 16 May – 17 June 2025.

SCAN HERE TO REGISTER



For any inquiries, feel free to contact rcc.sdgintl@gmail.com



I'll be doing the below for my system which I will be coding for the hackathon in the image above

DETAILED EXPLANATION

Project Overview: AI-Powered Holistic Health Companion

Core Vision: To establish a cutting-edge, user-centric digital platform that proactively empowers individuals in their health journey. This encompasses early detection of chronic conditions, intelligent health management, and personalized well-being recommendations, leveraging the latest advancements in AI and cloud computing.

Target Audiences: Individuals seeking proactive health management, those with potential diabetes risk, and anyone interested in personalized skincare and general health information.

Target Languages (User-Facing): Malay, English (Multilingual support for broad accessibility in Malaysia and beyond).

Overarching Technology Philosophy: Cloud-native, serverless-first (where appropriate), scalable, secure, and leveraging Google's AI ecosystem for advanced intelligence.

Feature-by-Feature Deep Dive: Technology, Implementation, and Interconnections

1. Voice-based Diabetes Risk Prediction

Description: This cornerstone feature leverages advanced acoustic analysis and machine learning to infer diabetes risk from a user's voice. The system records a brief voice sample, extracts a rich set of acoustic features (such as pitch variations, voice stability/jitter, energy contours, formants, and spectral characteristics), and feeds these into a trained predictive model. The output is a risk likelihood, not a diagnosis, intended for early awareness.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Non-communicable diseases, prevention & treatment): Directly addresses the need for non-invasive, accessible methods of early detection for non-communicable diseases like diabetes, which are a major global health burden. By providing an early indicator, it encourages timely medical consultation and lifestyle intervention.

Detailed Technologies & Implementation:

Frontend (React.js):

User Interface (UI): A clean, intuitive UI built with React components. This includes a prominent "Record Voice Sample" button, a real-time audio visualization (e.g., a waveform or volume meter) during recording for user feedback, and a progress indicator.

Audio Recording: Utilizes the browser's native MediaRecorder API. This API allows direct access to the user's microphone (with explicit user permission) and captures audio streams. The recorded audio is typically buffered as Blob objects.

Audio Pre-processing (Client-side): Before sending, the recorded audio Blob might be transcoded to a specific format (e.g., WAV or a high-quality MP3) using a client-side library (e.g., web-audio-api for basic processing or a specialized WASM-compiled audio encoder) to ensure compatibility and reduce payload size for the backend.

API Integration: Uses fetch or axios to send the processed audio data to the Node.js backend API endpoint.

Backend (Node.js - Express.js or similar framework):

API Endpoint: A dedicated RESTful API endpoint (/api/voice-analysis) designed to receive multipart/form-data containing the audio file.

Audio Ingestion & Validation: Receives the audio file. Initial validation checks (file type, size, duration) are performed.

Audio Persistence (Google Cloud Storage): The raw or pre-processed audio file is immediately uploaded to a designated bucket in Google Cloud Storage (GCS). GCS provides durable, scalable, and secure storage for large binary objects. This ensures data persistence and provides an accessible URI for the ML model.

Why GCS? It's highly scalable, cost-effective for large files, and integrates seamlessly with other Google Cloud services, especially Vertex AI.

Machine Learning (ML) Model Trigger: Once the audio is successfully stored in GCS, a trigger (e.g., a Firebase Cloud Function or a direct HTTP request from the Node.js backend) is invoked to initiate the ML inference process.

Feature Extraction (Backend/Vertex AI): While voice features (pitch, jitter, energy, MFCCs, etc.) could be extracted on the Node.js server using libraries like librosa (via a Python bridge) or specialized Node.js audio processing modules, for a production-grade ML pipeline, it's more common and efficient to have the feature extraction logic either:

Part of the deployed ML model: The model served on Vertex AI is designed to take the raw audio data (or a pre-processed stream) and perform the necessary feature extraction internally before prediction.

A separate pre-processing service: A small Firebase Cloud Function or a dedicated Cloud Run service could be responsible solely for taking the GCS audio URI, extracting features, and then passing these features to the actual prediction model.

Machine Learning Model Deployment (Google Cloud Platform):

Vertex AI (or Google Cloud AI Platform): This is the core ML operational environment. Your pre-trained machine learning model (e.g., built with TensorFlow, PyTorch, or Scikit-learn) is deployed here as a managed endpoint.

Model Training: The model would have been trained on a large, diverse dataset of voice samples from individuals with and without diabetes, carefully labeled. Feature engineering and selection (e.g., prosodic features, perturbation measures, spectral features) are critical. Common ML algorithms for such classification tasks include Support Vector Machines (SVMs), Random Forests, Gradient Boosting Machines (XGBoost/LightGBM), or even deep learning models like Convolutional Neural Networks (CNNs) or Recurrent Neural Networks (RNNs) if complex temporal patterns in voice are being analyzed.

Deployment: The model is deployed to a Vertex AI Endpoint, providing a stable, low-latency API for real-time predictions. The Node.js backend sends inference requests to this endpoint, passing either the raw audio (if the model handles extraction) or the extracted features.

Database (Firestore):

Data Storage: After receiving the prediction from Vertex AI, the Node.js backend stores the results in Firestore. Each document would typically contain: `userId`, `timestamp`, `audioFileRef` (GCS URI), `predictedRiskScore`, `confidenceScore`, `extractedFeatures` (if applicable), `rawDataProcessed` (boolean).

Why Firestore? It's a highly scalable, serverless NoSQL document database perfect for storing user-specific data, test results, and metadata, offering real-time synchronization capabilities for the dashboard.

2. Personal Health Tracker Dashboard

Description: This feature provides users with a centralized, interactive view of their health journey. It displays trends for voice-based risk predictions, allows manual logging and visualization of blood sugar readings, facilitates goal setting (e.g., weight loss targets, reduced sugar intake), and delivers timely alerts when adverse health trends are detected.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Non-communicable diseases, treatment & self-management): Promotes proactive disease management by enabling users to track key health indicators, set personal goals, and receive actionable insights, thereby fostering self-efficacy in health.

Detailed Technologies & Implementation:

Frontend (React.js):

Dashboard Layout: Responsive layout with distinct sections for different health metrics.

Data Visualization: Libraries like Chart.js or Recharts are used to render interactive line graphs (for trends over time), bar charts (for discrete measurements), and pie charts (for goal progress). These libraries efficiently consume JSON data fetched from the backend.

Form Management: React forms for inputting manual blood sugar readings, updating weight, and setting/modifying health goals. Input validation and user feedback are crucial here.

Real-time Updates: Leverages Firebase Firestore's real-time listeners. When data in Firestore changes (e.g., a new voice analysis result or a logged blood sugar reading), the React app can automatically update the dashboard without a full page refresh, providing a dynamic user experience.

Notification Display: UI components to display in-app notifications and alerts received from the backend (e.g., "Your diabetes risk trend is increasing!").

Backend (Node.js - Express.js / Firebase Cloud Functions):

API Endpoints: RESTful APIs (/api/health-data, /api/goals) for fetching and updating user health records and goals. These endpoints interact directly with Firebase Firestore.

Data Aggregation & Analysis (Firebase Cloud Functions):

Scheduled Functions: Cloud Functions can be scheduled to run periodically (e.g., daily, weekly) to analyze user data in Firestore.

Trend Analysis: Logic to calculate trends (e.g., moving averages of risk scores, rate of blood sugar increase). This involves querying Firestore for historical data and performing statistical calculations.

Alert Generation: Based on predefined thresholds or trend detection, Cloud Functions trigger alerts. These alerts can be stored in Firestore (for in-app display) or sent as push notifications.

Why Cloud Functions? They are serverless, cost-effective for event-driven or scheduled tasks, and scale automatically, perfect for background data processing and alert logic.

Database (Firebase Firestore):

Schema Design: Stores collections such as:

users: User profiles, general health info.

voiceTestResults: Linking back to feature 1, contains userId, timestamp, riskScore, etc.

bloodSugarReadings: userId, timestamp, readingValue, mealContext.

healthGoals: userId, goalType (e.g., weight_loss, sugar_reduction), targetValue, currentValue, startDate, endDate.

alerts: userId, timestamp, alertType, message, isRead.

Firestore's real-time capabilities are key here for responsive dashboards.

3. Doctor Advice Reliability Tracker

Description: This innovative feature aims to foster informed patient engagement and improve the quality of medical information. Users can input or transcribe advice received from their doctors. The system then leverages a sophisticated AI model to summarize this advice, cross-reference it against general medical knowledge, and even provide supplementary "doctor tips" or highlight areas for further clarification.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.c - Health workforce, efficiency & patient education): Empowers patients to better understand and critically evaluate medical advice, potentially leading to

improved adherence and health outcomes. It indirectly supports the health workforce by clarifying information.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Innovation in services, AI in healthcare): A prime example of applying advanced AI (LLMs) to enhance patient-doctor interactions and information management in healthcare.

Detailed Technologies & Implementation:

Frontend (React.js):

Input Interface: Text area for users to type or paste doctor's advice. Could also integrate a simple microphone input if transcribing short notes.

Feedback Mechanism: Buttons or a rating system for users to provide feedback on the advice (e.g., "helpful," "clear," "conflicting").

Display of AI Insights: Dedicated sections to show the Gemini-generated summary, reliability checks, and "doctor tips" clearly.

Backend (Node.js - Express.js / Firebase Cloud Functions):

API Endpoint: An API endpoint (/api/doctor-advice) to receive user-submitted advice.

AI Processing (Gemini API via Node.js Client Library):

Prompt Engineering: This is where the "experienced prompt engineer" aspect comes to life. Node.js backend constructs carefully designed prompts for Gemini.

For Summarization: "Summarize the following medical advice concisely, highlighting key recommendations and dosages: [User-provided advice text]"

For Cross-referencing/Verification: "Evaluate the following medical advice for a patient with diabetes against general medical guidelines. Is it consistent with standard care? Highlight any potential discrepancies or areas for further inquiry: [User-provided advice text]"

For "Doctor Tips": "Based on the following medical advice for diabetes management, generate 3 additional practical tips for the patient to consider: [User-provided advice text]"

API Calls: The Node.js backend makes secure API calls to the Gemini API endpoint, passing the crafted prompts and receiving the AI-generated responses.

Error Handling & Fallbacks: Robust error handling for API calls, including retries and informative messages if Gemini service is unavailable.

Database (Firebase Firestore):

Schema: doctorAdviceEntries collection with fields like userId, timestamp, originalAdviceText, geminiSummary, geminiReliabilityCheckOutput, geminiDoctorTips, userFeedbackRating.

Knowledge Base (Optional but Recommended): For the "cross-reference/verify" functionality, a curated, up-to-date knowledge base of diabetes management guidelines and common medical protocols could be stored in Firestore. This data would be retrieved by the Node.js backend and provided as context to Gemini in the prompt (Retrieval-Augmented Generation - RAG) to ensure higher accuracy and relevance.

Example RAG prompt: "Given the following medical guidelines for diabetes management: [Retrieved guidelines text]. Evaluate this doctor's advice: [User-provided advice text]. Is it consistent?"

4. User & Doctor Conversation Record & Summarization

Description: This advanced feature aims to capture and distill the essence of crucial patient-doctor dialogues. Users can record their conversations (with explicit consent from all parties). The system then transcribes the audio, generates concise summaries, identifies key medical terms, and can even provide "doctor tips" or double-check the advice discussed, serving as an invaluable memory aid and verification tool.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.c - Health workforce, patient communication & empowerment): Significantly improves patient comprehension, recall, and adherence to medical instructions, thereby boosting the effectiveness of healthcare interactions and patient empowerment.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Technological capabilities, AI in healthcare): A sophisticated application of multimodal AI (speech-to-text + LLM) to a critical healthcare process, enhancing information retention and accuracy.

Detailed Technologies & Implementation:

Frontend (React.js):

Audio Recording Interface: Similar to Feature 1, uses MediaRecorder API for capturing longer audio sessions. A clear visual indicator of recording status and duration.

Consent Management: A highly visible and mandatory consent prompt before recording, explaining data usage, privacy, and necessity for doctor's consent (for ethical and legal compliance).

Display of Summaries: Rich text display for conversation summaries, "doctor tips," and highlighted double-checked advice.

Playback Functionality: Simple audio playback for the recorded conversation.

Backend (Node.js - Express.js / Firebase Cloud Functions):

Audio Ingestion: API endpoint (/api/conversation-upload) to receive the potentially large audio files.

Secure Storage (Google Cloud Storage): Uploads raw audio files to a private GCS bucket. Crucial for privacy and security.

Audio Transcription (Google Cloud Platform):

Google Cloud Speech-to-Text API: This is the workhorse for converting spoken words to text. After the audio file is in GCS, a Firebase Cloud Function (triggered by the GCS upload event or an explicit call from Node.js) invokes the Speech-to-Text API.

Configuration: The API can be configured for long audio, speaker diarization (identifying different speakers, e.g., "Doctor said:", "Patient said:"), and language recognition (Malay, English).

Output: Returns a JSON object with transcribed text, timestamps for words, and speaker labels.

AI Processing (Gemini API via Node.js Client Library):

Prompt Engineering (Advanced): Once transcribed, the full conversation text (with speaker labels) is sent to Gemini. Prompts are crafted to perform multiple tasks:

Summarization: "Summarize the following doctor-patient conversation, highlighting key medical findings, diagnoses, treatment plans, and follow-up actions: [Transcribed conversation text]"

Key Term/Advice Extraction: "From the conversation, identify all specific medical advice given by the doctor, including medication names, dosages, lifestyle changes, and next steps: [Transcribed conversation text]"

"Doctor Tips" Generation: "Based on the summarized conversation and advice, generate 3 practical tips for the patient to enhance their understanding or adherence: [Summary/Extracted Advice]"

Advice Double-checking: "Cross-reference the following specific advice from the conversation with general medical guidelines for diabetes: [Specific advice snippet from conversation]. Indicate if it's consistent or requires further clarification."

Parallel/Sequential Calls: Depending on the complexity, multiple calls to Gemini might be made in parallel or sequence to get different types of analysis from the same conversation.

Database (Firebase Firestore):

Schema: conversations collection with fields like userId, timestamp, audioFileRef (GCS URI), transcribedText, geminiSummary, geminiExtractedAdvice, geminiDoctorTips, geminiCheckedAdviceResults.

Security Rules: Implement strict Firebase Security Rules to ensure only the authorized user can access their conversation data.

5. Health Chatbot (Multilingual)

Description: A simple, reliable conversational AI designed to provide answers to user questions about diabetes in both Malay and English. Crucially, the chatbot emphasizes factual accuracy and avoids AI "hallucinations" (generating plausible but incorrect information) by primarily relying on a verified knowledge base or a carefully fine-tuned model.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Health education & awareness): Improves access to reliable health information, especially on a prevalent non-communicable disease like diabetes, breaking down language barriers for wider public health literacy.

Detailed Technologies & Implementation:

Frontend (React.js):

Chat Interface: Standard chat UI with an input field for user queries and a scrollable message display area showing conversation history.

Language Selection: A dropdown or toggle for users to select their preferred language (Malay/English).

Real-time Updates: Displays bot responses as they are generated.

Backend (Node.js - Express.js / Firebase Cloud Functions):

API Endpoint: A dedicated endpoint (/api/chatbot) to receive user queries.

Query Processing: Receives the user's text query and selected language.

Multilingual Handling:

Input Language Detection (Optional but Recommended): If the user doesn't explicitly select a language, Google Cloud Translation API (or a smaller local library) can be used to detect the input language.

Language-Specific Knowledge Retrieval: If using a RAG approach, the backend retrieves relevant information from the Firestore knowledge base for the detected or selected language.

AI (Gemini API via Node.js Client Library):

Retrieval-Augmented Generation (RAG - Highly Recommended for Accuracy):

Knowledge Base (Firestore): A dedicated Firestore collection (e.g., diabetesKnowledge) stores well-structured, verified Q&A pairs, facts, and explanations about diabetes in both Malay and English. This is your curated truth source.

Retrieval: When a user asks a question, the Node.js backend performs a semantic search (using text embeddings or simple keyword matching on the Firestore knowledge base) to find the most relevant pieces of information.

Augmentation & Generation: The retrieved relevant information is then included in the prompt sent to Gemini: "Answer the following question in [language, e.g., Malay], strictly using the provided context. If the context does not contain the answer, state that you don't have enough information. Question: [User Query]. Context: [Retrieved knowledge base entries]."

Why RAG with Gemini? This architecture drastically reduces hallucinations. Gemini acts as a sophisticated text generator constrained by your verified data, rather than freely inventing information.

Fine-tuned Gemini Model (Alternative/Complementary): For specific, common questions, a smaller version of Gemini (if available for fine-tuning) could be fine-tuned on your diabetes Q&A dataset. This offers very fast, highly domain-specific responses without external retrieval. However, RAG is generally preferred for its auditability and ease of updating the knowledge base without retraining the model.

Database (Firestore):

diabetesKnowledge Collection: Stores structured health information in multiple languages. Example schema: id, question_en, answer_en, question_ms, answer_ms, keywords_en, keywords_ms, source_url.

chatbotConversations Collection: Stores chat history for user reference.

6. AI-Powered Skin Type Analysis & Skincare Recommendation

Description: This feature integrates computer vision with personalized recommendations. Users capture a selfie, which an AI model analyzes to classify their skin type (e.g., oily, dry, normal, combination, sensitive). Based on this classification, the system generates bespoke skincare routines, product suggestions, and relevant tips.

SDG Alignment:

SDG 3: Good Health and Well-being (Holistic Personal Care & Empowerment): Extends health beyond disease to include broader well-being, enabling users to make informed choices about personal care that can impact confidence and overall health.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Technological innovation in consumer health/beauty): Showcases advanced computer vision and AI applied to a consumer-facing health/beauty domain, driving innovation in personalized services.

Detailed Technologies & Implementation:

Frontend (React.js):

Camera Interface: Uses `MediaDevices.getUserMedia()` to display a live camera feed. Provides a "Capture Photo" button.

Image Pre-processing (Client-side): Before upload, the captured image (typically a Blob) is processed:

Resizing/Compression: Reduces image dimensions and quality to optimize upload speed and reduce storage/processing costs. Libraries like canvas or react-image-file-resizer can be used.

Cropping/Centering: Ensures the face is appropriately framed for the AI model.

Loading Indicators: Visual feedback during image upload and AI analysis.

Results Display: Clear display of detected skin type and a rich, interactive section for personalized skincare recommendations (e.g., product categories, ingredients to look for, routine steps).

Backend (Node.js - Express.js / Firebase Cloud Functions):

Image Upload API: An API endpoint (`/api/skin-analysis-upload`) designed to receive the image file.

Secure Image Storage (Google Cloud Storage): The uploaded image is stored in a private GCS bucket. Temporary storage is usually sufficient if analysis is real-time.

AI Model Trigger (Firebase Cloud Function): A Cloud Function triggered by the GCS upload event or an explicit call from Node.js invokes the computer vision model.

Computer Vision Model Deployment (Google Cloud Platform):

Vertex AI (Custom Models): This is where your custom image classification model resides.

Model Training: The model is trained on a massive dataset of facial images, meticulously labeled with corresponding skin types (e.g., using professional dermatological assessments). This involves collecting diverse images (different lighting, ethnicities, ages, skin conditions) and robust annotation.

Architecture: Typically, a Convolutional Neural Network (CNN) architecture (e.g., ResNet, EfficientNet) would be fine-tuned or trained from scratch for this image classification task.

Deployment: The trained model is deployed to a Vertex AI Prediction Endpoint, which provides a low-latency API for real-time inference. The Cloud Function or Node.js backend sends the GCS URI of the image (or the raw image data) to this endpoint.

Output: The model returns the predicted skin type (e.g., "Oily," "Dry," "Combination") and a confidence score.

Skincare Recommendation Logic (Node.js / Gemini API):

Skin Type to Prompt Mapping: Based on the skin type detected by Vertex AI, the Node.js backend constructs a precise prompt for Gemini.

AI (Gemini API):

Prompt Engineering: "Given a [Detected Skin Type, e.g., 'Oily'] skin type, generate a comprehensive daily and nightly skincare routine, including types of products (e.g., cleanser, serum, moisturizer), key ingredients to look for, and practical tips for management. Format the output as a step-by-step routine with bullet points for products/ingredients."

RAG for Specific Product Recommendations (Advanced): If you want to recommend specific products, you would extend this with a Firestore-based product catalog. Gemini would then be prompted to select suitable products from this catalog based on the detected skin type and the generated routine, explaining why those products are suitable.

Database (Firebase Firestore):

skinAnalysisResults Collection: Stores userId, timestamp, imageRef (GCS URI), detectedSkinType, confidence, geminiSkincareRecommendations.

skincareKnowledgeBase Collection (Optional but Recommended): Stores curated data on skincare ingredients, product types, routines, and tips linked to specific skin types. This would be used in a RAG setup for Gemini to ensure reliable recommendations.

This detailed breakdown provides a clear architectural vision for your project, highlighting the role of each technology in delivering a sophisticated and impactful solution. Remember, thorough testing, continuous model refinement, and adherence to privacy regulations (especially GDPR, PDPA in Malaysia) will be paramount for a successful and ethical deployment.

SHORT EXPLANATION

As an experienced prompt engineer, I can clearly outline your project's features, their SDG alignment, and the technologies used, emphasizing your commitment to cutting-edge solutions.

Your project is a holistic digital health companion, leveraging AI and multimodal capabilities to empower users in managing their well-being, from chronic disease risk assessment to personalized skincare.

Project Overview: AI-Powered Holistic Health Companion

Core Vision: To provide an intelligent, user-friendly platform for proactive health management, early disease detection, and personalized well-being advice.

Target Languages: Malay, English

Key Technologies: Google Cloud Platform (Vertex AI, Cloud Storage, Speech-to-Text, Cloud Functions), Firebase (Firestore, Authentication), Gemini API, React.js, Node.js

Features, SDG Alignment, and Technology Stack

Here's a detailed breakdown of each feature:

1. Voice-based Diabetes Risk Prediction

Description: This feature allows users to record a short voice sample (10-30 seconds), which is then analyzed by a custom machine learning model to predict the likelihood of diabetes or blood sugar irregularities. The analysis focuses on nuanced voice features like pitch, jitter, and energy.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Non-communicable diseases, prevention & treatment): Directly contributes to early detection and prevention of diabetes.

Technologies:

Frontend (React.js): MediaRecorder API or react-mic for audio capture; UI for recording and feedback.

Backend (Node.js): Audio processing libraries (e.g., web-audio-api on client for initial processing, server-side if needed).

Machine Learning (Google Cloud Platform):

Vertex AI (or Google Cloud AI Platform): For hosting and serving the custom pre-trained machine learning model (TensorFlow/PyTorch) for voice feature analysis and prediction.

Database (Firestore): Stores voice sample metadata, analysis results, and prediction scores.

Storage (Google Cloud Storage): Securely stores raw/processed audio files.

2. Personal Health Tracker Dashboard

Description: A comprehensive user dashboard for tracking health metrics over time. This includes visualizing voice-based test results, logging manual blood sugar readings, setting personalized health goals (e.g., weight loss, reduced sugar intake), and receiving proactive alerts for increasing diabetes risk trends.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Non-communicable diseases, prevention & treatment): Empowers users in proactive self-management and goal achievement related to health.

Technologies:

Frontend (React.js): Chart.js or Recharts for data visualization; forms for data input; UI for goal setting and alerts.

Backend (Node.js): API endpoints for data retrieval and storage; logic for risk trend analysis.

Serverless (Firebase Cloud Functions): Automated tasks for trend analysis and sending push notifications.

Database (Firestore): Stores all user health data, including voice test history, blood sugar logs, health goals, and notification preferences.

3. Doctor Advice Reliability Tracker

Description: This feature empowers users to track and evaluate the reliability and consistency of medical advice received from their healthcare providers. It allows for user input of advice, evaluation mechanisms, and cross-referencing with general medical guidelines.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.c - Health workforce, efficiency): Improves transparency and user engagement with medical advice, potentially leading to better adherence and outcomes.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Innovation in services): Application of AI to enhance healthcare information management.

Technologies:

Frontend (React.js): Interface for inputting advice; rating/feedback system.

Backend (Node.js): API endpoints for storing and retrieving advice entries.

AI (Gemini API):

Summarization: Condenses doctor's advice into concise summaries.

Cross-referencing/Verification: Evaluates advice against medical knowledge/guidelines using targeted prompts.

"Doctor Tips" Generation: Creates supplementary tips or questions based on summarized advice.

Database (Firebase Firestore): Stores user-entered advice, reliability ratings, and Gemini-generated insights.

4. User & Doctor Conversation Record & Summarization

Description: A powerful feature that allows users (with consent) to record conversations with their doctors. The system then transcribes and summarizes these discussions, enabling users to review key points, and provides "doctor tips" or double-checks the given medical advice.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.c - Health workforce, efficiency & patient empowerment): Improves patient understanding, recall, and adherence to medical advice, fostering better patient-doctor communication.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Technological capabilities): Innovative application of AI for healthcare information management.

Technologies:

Frontend (React.js): Audio recording interface; display for summaries and tips.

Backend (Node.js):

Audio Transcription (Google Cloud Platform):

Google Cloud Speech-to-Text API: High-accuracy transcription of spoken conversations into text.

AI (Gemini API):

Conversation Summarization: Extracts key points, decisions, and action items from transcripts.

Medical Term/Advice Identification: Highlights specific recommendations.

"Doctor Tips" Generation: Suggests follow-up questions, reminders, or general health tips.

Advice Double-checking: Cross-references specific advice against general medical knowledge.

Database (Firebase Firestore): Stores transcribed conversations, Gemini-generated summaries, and identified tips/checked advice.

Storage (Google Cloud Storage): Securely stores raw audio recordings of conversations.

5. Health Chatbot (Multilingual)

Description: A simple, reliable chatbot designed to answer user questions about diabetes in both Malay and English. It prioritizes accuracy and avoids hallucinations by utilizing a controlled knowledge base or a fine-tuned small language model.

SDG Alignment:

SDG 3: Good Health and Well-being (Target 3.4 - Health education & awareness): Provides accessible, reliable health information, contributing to public health literacy.

Technologies:

Frontend (React.js): Chat interface for user interaction.

Backend (Node.js): API endpoint for chatbot interactions.

AI (Gemini API):

Fine-tuned Gemini Model: Can be fine-tuned on a curated dataset of diabetes Q&A in Malay and English for controlled responses.

Retrieval-Augmented Generation (RAG) with Gemini: Combines a verified knowledge base (from Firestore) with Gemini to formulate accurate answers, minimizing hallucinations.

Database (Firebase Firestore): Stores the chatbot's knowledge base (Q&A pairs, diabetes facts).

6. AI-Powered Skin Type Analysis & Skincare Recommendation

Description: This feature enables users to scan their face via camera. An AI model analyzes the image to accurately detect the user's skin type (oily, normal, dry, combination). Based on this detection, the system then provides personalized skincare recommendations, including suitable products, routines, and tips.

SDG Alignment:

SDG 3: Good Health and Well-being (Broader well-being and self-care): Promotes holistic personal care and empowers users with personalized beauty and health advice.

SDG 9: Industry, Innovation and Infrastructure (Target 9.5 - Technological innovation in consumer services): Exemplifies advanced technological application (computer vision, AI) in personal health and beauty.

Technologies:

Frontend (React.js): Camera access (`MediaDevices.getUserMedia()`); image capture; client-side image pre-processing; UI for displaying results and recommendations.

Backend (Node.js): API endpoint to receive images.

Image Storage (Google Cloud Storage): Securely stores captured facial images.

Computer Vision (Google Cloud Platform):

Vertex AI (Custom Models): For training and deploying a custom image classification model specifically for skin type detection from facial images.

Vertex AI Prediction Endpoint: To serve the trained model for inference.

AI (Gemini API):

Recommendation Generation: Crafts personalized skincare routines, product suggestions, and tips based on the detected skin type (leveraging prompt engineering and potentially a RAG approach with a Firestore-based skincare knowledge base).

Database (Firebase Firestore): Stores user's skin analysis history and the curated skincare knowledge base.

Your project masterfully integrates advanced AI capabilities, cloud infrastructure, and a user-centric design to deliver a comprehensive health and wellness solution. This demonstrates a strong alignment with critical Sustainable Development Goals by fostering health, well-being, and technological innovation.